

1_MLBasics

March 16, 2025

0.1 CS-E4740 - Federated Learning D (Spring 25)

1 Assignment 1: ML Basics

1.0.1 R. Gafur, A. Jung

<h2>Deadline: 17.03.2025</h2>

1.1 Learning Goals

Access weather data from the Finnish Meteorological Institute (FMI).

Utilize Python libraries (scikit-learn, Keras) to train basic machine learning (ML) models.

Implement regularization techniques via data augmentation.

1.2 Background Material

- Chapter 2 of [FLBook \(PDF\)](#)
- optional: [Linear model implementation in scikitlearn](#), [Decision tree implementation in scikitlearn](#), [Convolutional Neural Networks \(CNN\) implementation in Keras](#)

1.3 Loading the Data

Before performing any data preprocessing or model training, we begin by loading the dataset and conducting an initial inspection.

Steps:

1. **Load the dataset:** We read the weather data from a CSV file using `pandas`.
2. **Check data types:** Displaying column data types helps us understand the structure of the dataset.
3. **Inspect the first data point:** We print all feature values of the first record to get an overview of the dataset content.
4. **Check dataset dimensions:** The dataset shape is printed to determine the number of rows and columns.

```
[2]: # General libraries
import pandas as pd
import numpy as np
```

```

# Tensorflow
import tensorflow as tf # for CNN model (tf.keras)

# Scikit-learn
from sklearn.tree import DecisionTreeRegressor # for Decision Tree Regressor
from sklearn.linear_model import LinearRegression # for Linear Regressioin
from sklearn.metrics import mean_squared_error # to measure mean squared error
↳ (MSE)
from sklearn.preprocessing import StandardScaler # to standardize the features

# Load the data
dataset = pd.read_csv('weather_data.csv')

# Check the data types and
print(dataset.dtypes)
print("\n*****\n")
print("First data point:")
print(dataset.iloc[0])
print("\n*****\n")
print(f"Dataset Shape: {dataset.shape}")

```

```

Observation station      object
Year                     int64
Month                    int64
Day                      int64
Time [Local time]        object
Average temperature [°C] float64
Maximum temperature [°C] float64
Minimum temperature [°C] float64
Average relative humidity [%] int64
Wind speed [m/s]         object
Maximum wind speed [m/s] object
Average wind direction [°] object
Maximum gust speed [m/s] object
Precipitation [mm]       float64
Average air pressure [hPa] float64
dtype: object

```

```

First data point:
Observation station      Kustavi Isokari
Year                     2023
Month                    4
Day                      1
Time [Local time]        00:00
Average temperature [°C] -0.6
Maximum temperature [°C] -0.4

```

Minimum temperature [°C]	-0.8
Average relative humidity [%]	82
Wind speed [m/s]	6
Maximum wind speed [m/s]	6.4
Average wind direction [°]	15
Maximum gust speed [m/s]	7.6
Precipitation [mm]	0.0
Average air pressure [hPa]	1007.5

Name: 0, dtype: object

Dataset Shape: (720, 15)

1.3.1 Turning raw data into datapoints, characterized by features and label.

The dataset consists of hourly weather measurements, including:

- **Temperature:** Average, Maximum, Minimum ([°C])
- **Humidity:** Average relative humidity ([%])
- **Wind:** Speed, Maximum speed, Average direction ([°]), Maximum gust speed ([m/s])
- **Precipitation:** Total precipitation ([mm])
- **Air Pressure:** Average air pressure ([hPa])

1.3.2 Data Points

A data point corresponds to a specific hour, e.g., *06-April-2023 from 06:00 - 07:00*, which is recorded as 2023-04-06 06:00:00 (starting time) after preprocessing.

- **Features** include all hourly observations from the **previous 5 hours**. Example: For the data point 2023-04-06 06:00:00, the features correspond to data from *01:00 - 06:00* on the same day.
- **Label** is the **temperature recorded 5 hours ahead**. Example: For the data point 2023-04-06 06:00:00, the label corresponds to the temperature measured during *11:00 - 12:00*.
- **Dataset Splits:**
 - **Training Set:** Includes data from 2023-04-06 00:00:00 to 2023-04-08 00:00:00.
 - **Validation Set:** Comprises the remaining hours of 2023.

```
[3]: # Copy the main dataset
data = dataset.copy()

# Convert specified columns to numeric (handling errors with 'coerce')
numeric_columns = [
    'Wind speed [m/s]',
    'Maximum wind speed [m/s]',
    'Average wind direction [°]',
```

```

        'Maximum gust speed [m/s]'
    ]
    data[numeric_columns] = data[numeric_columns].apply(pd.to_numeric,
        errors='coerce')

    # Fill missing values with 0
    data.fillna(0, inplace=True)

    # Create a 'timestamp' column by combining year, month, day, and local time
    data['timestamp'] = pd.to_datetime(
        data['Year'].astype(str) + '-' +
        data['Month'].astype(str) + '-' +
        data['Day'].astype(str) + ' ' +
        data['Time [Local time]']
    )

    # Identify columns for lagged features (excluding non-relevant ones)
    excluded_columns = ['Observation station', 'timestamp', 'Year', 'Month', 'Day',
        'Time [Local time]']
    columns_to_lag = [col for col in data.columns if col not in excluded_columns]

    # Create lagged features for the previous 5 hours
    for lag in range(1, 6):
        for col in columns_to_lag:
            data[f"{col} lag_{lag}"] = data[col].shift(lag)

    # Define target variable (y) as the average temperature 5 hours ahead
    data['y'] = data['Average temperature [°C]'].shift(-5)

    # Drop rows with NaN values caused by shifts
    data.dropna(inplace=True)

    # Print dataset shape
    print(f>Data shape: {data.shape}")

```

Data shape: (710, 67)

```

[4]: # Sanity checks

    # Check the data rows number
    assert data.shape[0] == 710, f"Unexpected number of rows: {data.shape[0]}."

    print('Sanity check passed!')

```

Sanity check passed!

1.3.3 TASK 1.1: Split the Data into Training and Validation Sets

Task Description:

1. **Split the dataset** into training (X_{train} , y_{train}) and validation (X_{val} , y_{val}) sets based on the specified time range.
2. **Training Set:** Includes data from hours 2023-04-06 00:00:00 to (and including) 2023-04-08 00:00:00.
3. **Validation Set:** Consists of the remaining data.

Hints:

- Use **lagged feature columns** (i.e., columns containing "lag") to create feature sets.
- Assign **lagged columns** as the feature variables for X_{train} and X_{val} .
- Use the **y column** as the target variable for y_{train} and y_{val} .

Points: 0.5

```
[5]: train_start = '2023-04-06 00:00:00'
train_end = '2023-04-08 00:00:00'

train_data = data[(data['timestamp'] >= train_start) & (data['timestamp'] <=
    ↪train_end)]
val_data = data[data['timestamp'] > train_end]

# Select feature columns (lagged features only)
feature_columns = [col for col in data.columns if "lag" in col]

# Split the data into training and validation sets
# Split into features (X) and target (y)
X_train = train_data[feature_columns]
y_train = train_data['y']
X_val = val_data[feature_columns]
y_val = val_data['y']

print(f"X_train: {X_train.shape}, y_train: {y_train.shape}, X_val: {X_val.
    ↪shape}, y_val: {y_val.shape}")
```

X_train: (49, 50), y_train: (49,), X_val: (546, 50), y_val: (546,)

```
[6]: # Sanity checks

# Check the data rows number
assert X_train.shape[0] == 49, f"Unexpected number of rows for X_train:
    ↪{X_train.shape[0]}."
assert X_val.shape[0] == 546, f"Unexpected number of rows for X_val: {X_val.
    ↪shape[0]}."

print('Sanity check passed!')
```

Sanity check passed!

1.3.4 Preparing Data for Model Training: Standardization & Reproducibility

After splitting the dataset into training and validation sets, the next step is to **prepare the features** for model training. Many machine learning models, especially those based on gradient-based optimization (e.g., linear regression, neural networks), perform better when features are standardized.

Why Standardization? Feature standardization ensures that all input variables: - Have a **mean of 0** and **standard deviation of 1**, preventing features with different scales from disproportionately influencing the model. - Enable **faster and more stable convergence** during training.

Steps:

1. Standardize the Features:

- The `StandardScaler` from `scikit-learn` is used to transform both the training and validation datasets.
- `fit_transform()` is applied to `X_train` to compute and apply the transformation.
- `transform()` is applied to `X_val` to use the same scaling parameters as `X_train`, ensuring consistency.

2. Ensure Reproducibility:

- To achieve consistent results across different runs, we set random seeds for both **NumPy** and **TensorFlow**.

```
[8]: # Import the required library for feature scaling
from sklearn.preprocessing import StandardScaler

# Initialize the StandardScaler
scaler = StandardScaler()

# Fit the scaler to the training data and transform it
# This ensures that the mean is 0 and the standard deviation is 1 for each
# feature in X_train
X_train_scaled = scaler.fit_transform(X_train)

# Apply the same transformation to the validation set
# We use transform() instead of fit_transform() to ensure the same scaling
# parameters from X_train are applied
X_val_scaled = scaler.transform(X_val)

# Set random seed for NumPy to ensure reproducibility in any random operations
np.random.seed(42)

# Set random seed for TensorFlow to ensure reproducibility in model training
# and initialization
tf.random.set_seed(42)
```

1.3.5 TASK 1.2: Linear Regression

Task Overview: Now that we have **standardized** our features to ensure proper scaling, we can proceed with training our first machine learning model. In this task, you will implement and evaluate a **linear regression model** using the standardized training and validation datasets.

Task Instructions:

1. **Train a Linear Regression Model** using `X_train_scaled` as input features.
2. **Evaluate Model Performance** by computing the average squared error of the trained model on both the training (`X_train_scaled, y_train`) and validation (`X_val_scaled, y_val`) sets.

Points: 1.5

```
[9]: # Initialize and train the Linear Regression model
regressor = LinearRegression()
regressor.fit(X_train_scaled, y_train)

# Predict on the training and validation sets
y_train_pred = regressor.predict(X_train_scaled)
y_val_pred = regressor.predict(X_val_scaled)

# Compute Mean Squared Error (MSE) for training and validation sets
reg_train_error = mean_squared_error(y_train, y_train_pred)
reg_val_error = mean_squared_error(y_val, y_val_pred)

# Display results
print("\n***** Linear Regression Diagnosis *****")
print(f"Training Error (MSE): {reg_train_error:.4f}")
print(f"Validation Error (MSE): {reg_val_error:.4f}")
```

```
***** Linear Regression Diagnosis *****
Training Error (MSE): 0.0000
Validation Error (MSE): 22.3691
```

```
[10]: # Sanity Checks: Ensuring Correctness of Model Evaluation

# Check if computed errors are numeric values
assert isinstance(reg_train_error, float), "Error: reg_train_error must be a_
↳numeric value."
assert isinstance(reg_val_error, float), "Error: reg_val_error must be a_
↳numeric value."

print("Sanity check passed! The computed errors are valid numerical values.")
```

Sanity check passed! The computed errors are valid numerical values.

1.3.6 TASK 1.3: Convolutional Neural Network (CNN)

Task Overview: Now that we have trained a **linear regression model**, let's explore a more powerful approach: a **1D Convolutional Neural Network (CNN)**. CNNs are well-suited for sequential data as they can capture local patterns and temporal dependencies.

Task Instructions:

1. **Implement a 1D CNN Model** using `tf.keras.Sequential` with `X_train_scaled` as input.
2. **Define the CNN architecture** with the following layers:
 - **Input Layer:** Shape `(X_train_scaled.shape[1], 1)`.
 - **Multiple Conv1D Layers** with decreasing filter sizes and ReLU activation.
 - **Flatten Layer** to convert feature maps into a vector.
 - **Dense Output Layer** with a single neuron for regression.
3. **Compile the Model** with:
 - **Optimizer:** Adam
 - **Loss Function:** Mean Squared Error (MSE)
4. **Train the Model** on the training set.
5. **Evaluate Model Performance** on both the training and validation sets (similar to Task 1.2).

Hints:

- Use `tf.keras.layers.Conv1D()` to define the convolutional layers with appropriate filters, kernel_size, and activation.
- Use `tf.keras.layers.Flatten()` to reshape the convolutional output before passing it to the dense layer.
- Use `tf.keras.layers.Dense()` for the final output layer.
- Construct the model using `tf.keras.Sequential()`, and compile it with `model.compile()`.

This CNN will serve as a more expressive alternative to linear regression, allowing us to compare their performances.

Points: 2

```
[11]: # Reshape input data to match Conv1D expected shape (samples, timesteps,
      ↪ features)
X_train_cnn = X_train_scaled.reshape(X_train_scaled.shape[0], X_train_scaled.
      ↪ shape[1], 1)
X_val_cnn = X_val_scaled.reshape(X_val_scaled.shape[0], X_val_scaled.shape[1],
      ↪ 1)

# Define the CNN model
cnn_model = tf.keras.Sequential([
    tf.keras.layers.Conv1D(filters=64, kernel_size=3, activation='relu',
      ↪ input_shape=(X_train_cnn.shape[1], 1)),
```



```

tf.keras.layers.Conv1D(filters=32, kernel_size=3, activation='relu'),
tf.keras.layers.Conv1D(filters=16, kernel_size=3, activation='relu'),
tf.keras.layers.Flatten(),
tf.keras.layers.Dense(16, activation='relu'),
tf.keras.layers.Dense(1) # Output layer for regression
])

# Compile the model
cnn_model.compile(optimizer='adam', loss='mse')

# Train the model
cnn_history = cnn_model.fit(X_train_cnn, y_train, epochs=50, batch_size=16,
    ↪ validation_data=(X_val_cnn, y_val), verbose=1)

# Evaluate the model
cnn_train_error = cnn_model.evaluate(X_train_cnn, y_train, verbose=0)
cnn_val_error = cnn_model.evaluate(X_val_cnn, y_val, verbose=0)

print("\n***** Convolutional Neural Network (CNN) Diagnosis_
    ↪ *****")
print("Training error:", cnn_train_error)
print("Validation error:", cnn_val_error)

```

Epoch 1/50

/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
4/4          1s 28ms/step - loss:
1.9546 - val_loss: 8.1712
```

Epoch 2/50

```
4/4          0s 11ms/step - loss:
0.7802 - val_loss: 20.2136
```

Epoch 3/50

```
4/4          0s 12ms/step - loss:
0.2707 - val_loss: 34.0108
```

Epoch 4/50

```
4/4          0s 11ms/step - loss:
0.3195 - val_loss: 12.2149
```

Epoch 5/50

```
4/4          0s 19ms/step - loss:
0.1175 - val_loss: 7.7711
```

Epoch 6/50

```
4/4          0s 11ms/step - loss:
```

0.1627 - val_loss: 8.3098
Epoch 7/50
4/4 0s 12ms/step - loss:
0.1093 - val_loss: 11.4444
Epoch 8/50
4/4 0s 12ms/step - loss:
0.1036 - val_loss: 12.9782
Epoch 9/50
4/4 0s 12ms/step - loss:
0.1136 - val_loss: 10.0023
Epoch 10/50
4/4 0s 11ms/step - loss:
0.0828 - val_loss: 8.0377
Epoch 11/50
4/4 0s 12ms/step - loss:
0.0812 - val_loss: 7.9834
Epoch 12/50
4/4 0s 12ms/step - loss:
0.0721 - val_loss: 8.8459
Epoch 13/50
4/4 0s 12ms/step - loss:
0.0717 - val_loss: 8.8929
Epoch 14/50
4/4 0s 12ms/step - loss:
0.0693 - val_loss: 7.9264
Epoch 15/50
4/4 0s 12ms/step - loss:
0.0598 - val_loss: 7.4091
Epoch 16/50
4/4 0s 12ms/step - loss:
0.0554 - val_loss: 7.5467
Epoch 17/50
4/4 0s 12ms/step - loss:
0.0517 - val_loss: 7.7742
Epoch 18/50
4/4 0s 12ms/step - loss:
0.0504 - val_loss: 7.5276
Epoch 19/50
4/4 0s 12ms/step - loss:
0.0467 - val_loss: 7.1900
Epoch 20/50
4/4 0s 12ms/step - loss:
0.0433 - val_loss: 7.1337
Epoch 21/50
4/4 0s 12ms/step - loss:
0.0406 - val_loss: 7.1680
Epoch 22/50
4/4 0s 12ms/step - loss:

0.0386 - val_loss: 7.0438
Epoch 23/50
4/4 0s 11ms/step - loss:
0.0359 - val_loss: 6.9016
Epoch 24/50
4/4 0s 12ms/step - loss:
0.0333 - val_loss: 6.8995
Epoch 25/50
4/4 0s 12ms/step - loss:
0.0313 - val_loss: 6.9450
Epoch 26/50
4/4 0s 12ms/step - loss:
0.0296 - val_loss: 6.9394
Epoch 27/50
4/4 0s 12ms/step - loss:
0.0276 - val_loss: 6.9361
Epoch 28/50
4/4 0s 12ms/step - loss:
0.0258 - val_loss: 6.9753
Epoch 29/50
4/4 0s 12ms/step - loss:
0.0241 - val_loss: 7.0047
Epoch 30/50
4/4 0s 12ms/step - loss:
0.0225 - val_loss: 6.9994
Epoch 31/50
4/4 0s 12ms/step - loss:
0.0210 - val_loss: 7.0136
Epoch 32/50
4/4 0s 12ms/step - loss:
0.0196 - val_loss: 7.0507
Epoch 33/50
4/4 0s 21ms/step - loss:
0.0183 - val_loss: 7.0651
Epoch 34/50
4/4 0s 12ms/step - loss:
0.0171 - val_loss: 7.0636
Epoch 35/50
4/4 0s 12ms/step - loss:
0.0159 - val_loss: 7.0860
Epoch 36/50
4/4 0s 12ms/step - loss:
0.0149 - val_loss: 7.1105
Epoch 37/50
4/4 0s 12ms/step - loss:
0.0139 - val_loss: 7.1185
Epoch 38/50
4/4 0s 12ms/step - loss:

```

0.0130 - val_loss: 7.1361
Epoch 39/50
4/4          0s 12ms/step - loss:
0.0122 - val_loss: 7.1451
Epoch 40/50
4/4          0s 12ms/step - loss:
0.0114 - val_loss: 7.1447
Epoch 41/50
4/4          0s 12ms/step - loss:
0.0106 - val_loss: 7.1511
Epoch 42/50
4/4          0s 12ms/step - loss:
0.0099 - val_loss: 7.1428
Epoch 43/50
4/4          0s 12ms/step - loss:
0.0093 - val_loss: 7.1265
Epoch 44/50
4/4          0s 12ms/step - loss:
0.0087 - val_loss: 7.1272
Epoch 45/50
4/4          0s 12ms/step - loss:
0.0081 - val_loss: 7.1377
Epoch 46/50
4/4          0s 12ms/step - loss:
0.0076 - val_loss: 7.1437
Epoch 47/50
4/4          0s 15ms/step - loss:
0.0071 - val_loss: 7.1623
Epoch 48/50
4/4          0s 12ms/step - loss:
0.0066 - val_loss: 7.1919
Epoch 49/50
4/4          0s 12ms/step - loss:
0.0062 - val_loss: 7.2208
Epoch 50/50
4/4          0s 12ms/step - loss:
0.0058 - val_loss: 7.2421

```

***** Convolutional Neural Network (CNN) Diagnosis *****

Training error: 0.0049854451790452

Validation error: 7.242134094238281

```

[12]: # Sanity checks

# Check if the variables store numeric values
assert isinstance(cnn_train_error, float), "cnn_train_error must be a numeric_
↳value."

```

```

assert isinstance(cnn_val_error, float), "cnn_val_error must be a numeric value.
↪"

print('Sanity check passed!')

```

Sanity check passed!

1.3.7 TASK 1.4: Decision Tree Regressor

Task Overview: Building on our previous models, we will now implement a **Decision Tree Regressor**, a non-parametric model that can capture nonlinear relationships in the data.

Task Instructions:

1. **Train a Decision Tree Regressor** using `X_train_scaled` as input features with `max_depth=3`.
2. **Compute Mean Squared Error (MSE)** for training (`X_train_scaled, y_train`) and validation (`X_val_scaled, y_val`) sets.

Points: 1.5

```

[13]: # Initialize and train the Decision Tree Regressor with max_depth=3
tree_regressor = DecisionTreeRegressor(max_depth=3)
tree_regressor.fit(X_train_scaled, y_train)

# Predict on the training and validation sets
y_train_pred_tree = tree_regressor.predict(X_train_scaled)
y_val_pred_tree = tree_regressor.predict(X_val_scaled)

# Compute Mean Squared Error (MSE) for training and validation sets
tree_train_error = mean_squared_error(y_train, y_train_pred_tree)
tree_val_error = mean_squared_error(y_val, y_val_pred_tree)

print("\n***** Decision Tree Regressor Diagnosis *****")
print("Training error:", tree_train_error)
print("Validation error:", tree_val_error)

```

```

***** Decision Tree Regressor Diagnosis *****
Training error: 0.023813573048266926
Validation error: 12.172040647031434

```

```

[14]: # Sanity checks

# Check if the variables store numeric values
assert isinstance(tree_train_error, float), "tree_train_error must be a numeric_
↪value."
assert isinstance(tree_val_error, float), "tree_val_error must be a numeric_
↪value."

```

```
print('Sanity check passed!')
```

Sanity check passed!

1.4 Regularization

So far, we have trained **three basic ML models**, a linear model, a CNN and a decision tree. While these models can effectively learn patterns from the training data, they are **prone to overfitting**, especially when the dataset is small. To improve generalization, we can use **regularization techniques** such as adding the **ridge penalty**.

1.4.1 Implementing the Ridge Penalty with Data Augmentation

To implement **ridge regularization**, we will: 1. **Generate an augmented training set** using independent and identically distributed (i.i.d.) **Gaussian random vectors**. 2. **Set the labels** for this augmented data to **0**. 3. **Modify our existing models** (Linear Regression, CNN, and Decision Tree) to incorporate this augmented dataset.

This approach introduces a **penalty term**, encouraging models to keep their predictions closer to zero for these artificial data points, thus improving stability and reducing overfitting.

```
[15]: import numpy as np
import pandas as pd
import tensorflow as tf

def generate_augmented_dataset(X_train, y_train, augmentation_ratio=0.5,
    random_seed=42):
    """
    Generates an augmented training set by adding synthetic data points drawn
    from a Gaussian distribution.

    Parameters:
    - X_train (numpy array or DataFrame): Original training feature set.
    - y_train (numpy array or Series): Original training labels.
    - augmentation_ratio (float): Ratio of synthetic samples to real training
    samples (default: 0.5).
    - random_seed (int): Random seed for reproducibility.

    Returns:
    - X_train_aug (numpy array): Augmented feature set.
    - y_train_aug (numpy array): Augmented labels (including original and
    synthetic samples).
    """

    # Set random seed for full reproducibility
    np.random.seed(random_seed)
    tf.random.set_seed(random_seed) # Ensure reproducibility in TensorFlow if
    used later
```

```

# Determine the number of augmented samples
num_augmented_samples = int(augmentation_ratio * X_train.shape[0])

# Compute mean and standard deviation of each feature in X_train
feature_means = np.mean(X_train, axis=0)
feature_stds = np.std(X_train, axis=0) + 1e-8 # Small epsilon to avoid
↳ zero variance issues

# Generate synthetic feature vectors from a Gaussian distribution
X_augmented = np.random.normal(loc=feature_means, scale=feature_stds,
↳ size=(num_augmented_samples, X_train.shape[1]))

# Set the labels for the augmented data points to 0
y_augmented = np.zeros(num_augmented_samples)

# Combine the original and augmented datasets
X_train_aug = np.vstack((X_train, X_augmented))
y_train_aug = np.hstack((y_train, y_augmented))

return X_train_aug, y_train_aug

# Example usage:
# Assuming X_train_scaled and y_train are already defined
X_train_aug, y_train_aug = generate_augmented_dataset(X_train_scaled, y_train)

print("Original training set shape:", X_train_scaled.shape)
print("Augmented training set shape:", X_train_aug.shape)

```

```

Original training set shape: (49, 50)
Augmented training set shape: (73, 50)

```

1.4.2 TASK 1.5: Regularized Model Training with Augmented Data

In this task, you will **regularize the training** of all three models from previous tasks—**Linear Regression, CNN, and Decision Tree Regressor**—by replacing the original training set with the **augmented training set** generated above. The augmented dataset contains **synthetic feature vectors drawn from a Gaussian distribution**, with their labels set to **0**. For parametric models, using this augmented dataset to train ML a model is equivalent to adding a ridge penalty term (see Sec. 6.6. of [MLBook \(PDF\)](#)).

Task Instructions:

1. Train the following models using the augmented training set (`X_train_aug`, `y_train_aug`):
 - Linear Regression
 - Convolutional Neural Network (CNN)
 - Decision Tree Regressor

2. Compute Mean Squared Error (MSE) of these regularized models on the original training (X_train_scaled,y_train) and validation sets (X_val_scaled,y_val).

Points: 1.5

```
[18]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Conv1D, Flatten, Dense

      # Train a Linear Regression model using the augmented dataset
      custom_reg = LinearRegression()
      custom_reg.fit(X_train_aug, y_train_aug)

      # Train a CNN using the augmented dataset
      # Reshape X_train_aug for Conv1D
      X_train_cnn_aug = X_train_aug.reshape(X_train_aug.shape[0], X_train_aug.
      ↪shape[1], 1)

      # Define the CNN model
      custom_cnn = Sequential([
          Conv1D(filters=64, kernel_size=3, activation='relu',
          ↪input_shape=(X_train_cnn_aug.shape[1], 1)),
          Conv1D(filters=32, kernel_size=3, activation='relu'),
          Conv1D(filters=16, kernel_size=3, activation='relu'),
          Flatten(),
          Dense(16, activation='relu'),
          Dense(1) # Output layer for regression
      ])

      # Compile and train the CNN model
      custom_cnn.compile(optimizer='adam', loss='mse')
      custom_cnn.fit(X_train_cnn_aug, y_train_aug, epochs=50, batch_size=16,
      ↪verbose=1)

      # Train a Decision Tree Regressor using the augmented dataset
      custom_tree = DecisionTreeRegressor(max_depth=3)
      custom_tree.fit(X_train_aug, y_train_aug)

      custom_reg_train_error = mean_squared_error(y_train, custom_reg.
      ↪predict(X_train_scaled))
      custom_reg_val_error = mean_squared_error(y_val, custom_reg.
      ↪predict(X_val_scaled))

      print("\n***** Diagnosis of Regularized Linear Model *****")
      print("Training error:", custom_reg_train_error)
      print("Validation error:", custom_reg_val_error)
```



```

custom_cnn_train_error = custom_cnn.evaluate(X_train_scaled[...], np.newaxis],  

    ↪y_train, verbose=0)  

custom_cnn_val_error = custom_cnn.evaluate(X_val_scaled[...], np.newaxis],  

    ↪y_val, verbose=0)  
  

print("\n***** Diagnosis of Regularized CNN *****")  

print("Training error:", custom_cnn_train_error)  

print("Validation error:", custom_cnn_val_error)  
  

custom_tree_train_error = mean_squared_error(y_train, custom_tree.  

    ↪predict(X_train_scaled))  

custom_tree_val_error = mean_squared_error(y_val, custom_tree.  

    ↪predict(X_val_scaled))  
  

print("\n***** Diagnosis of Regularized Decision Tree  

    ↪*****")  

print("Training error:", custom_tree_train_error)  

print("Validation error:", custom_tree_val_error)

```

Epoch 1/50

/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

5/5 1s 3ms/step - loss:

1.6438

Epoch 2/50

5/5 0s 3ms/step - loss:

0.7619

Epoch 3/50

5/5 0s 3ms/step - loss:

0.5498

Epoch 4/50

5/5 0s 3ms/step - loss:

0.4007

Epoch 5/50

5/5 0s 3ms/step - loss:

0.3484

Epoch 6/50

5/5 0s 3ms/step - loss:

0.2927

Epoch 7/50

5/5 0s 2ms/step - loss:

0.2172

Epoch 8/50
5/5 0s 2ms/step - loss:
0.1588

Epoch 9/50
5/5 0s 10ms/step - loss:
0.1161

Epoch 10/50
5/5 0s 3ms/step - loss:
0.0822

Epoch 11/50
5/5 0s 3ms/step - loss:
0.0641

Epoch 12/50
5/5 0s 3ms/step - loss:
0.0503

Epoch 13/50
5/5 0s 3ms/step - loss:
0.0389

Epoch 14/50
5/5 0s 3ms/step - loss:
0.0309

Epoch 15/50
5/5 0s 2ms/step - loss:
0.0249

Epoch 16/50
5/5 0s 2ms/step - loss:
0.0202

Epoch 17/50
5/5 0s 3ms/step - loss:
0.0164

Epoch 18/50
5/5 0s 3ms/step - loss:
0.0138

Epoch 19/50
5/5 0s 3ms/step - loss:
0.0120

Epoch 20/50
5/5 0s 3ms/step - loss:
0.0102

Epoch 21/50
5/5 0s 3ms/step - loss:
0.0090

Epoch 22/50
5/5 0s 3ms/step - loss:
0.0080

Epoch 23/50
5/5 0s 3ms/step - loss:
0.0071

```

Epoch 24/50
5/5      0s 3ms/step - loss:
0.0063
Epoch 25/50
5/5      0s 3ms/step - loss:
0.0056
Epoch 26/50
5/5      0s 3ms/step - loss:
0.0051
Epoch 27/50
5/5      0s 3ms/step - loss:
0.0045
Epoch 28/50
5/5      0s 3ms/step - loss:
0.0041
Epoch 29/50
5/5      0s 3ms/step - loss:
0.0037
Epoch 30/50
5/5      0s 3ms/step - loss:
0.0033
Epoch 31/50
5/5      0s 3ms/step - loss:
0.0029
Epoch 32/50
5/5      0s 3ms/step - loss:
0.0026
Epoch 33/50
5/5      0s 3ms/step - loss:
0.0023
Epoch 34/50
5/5      0s 3ms/step - loss:
0.0020
Epoch 35/50
5/5      0s 3ms/step - loss:
0.0018
Epoch 36/50
5/5      0s 3ms/step - loss:
0.0016
Epoch 37/50
5/5      0s 10ms/step - loss:
0.0014
Epoch 38/50
5/5      0s 3ms/step - loss:
0.0012
Epoch 39/50
5/5      0s 3ms/step - loss:
0.0011

```

```

Epoch 40/50
5/5          0s 3ms/step - loss:
9.8457e-04
Epoch 41/50
5/5          0s 3ms/step - loss:
8.7058e-04
Epoch 42/50
5/5          0s 3ms/step - loss:
7.7667e-04
Epoch 43/50
5/5          0s 3ms/step - loss:
6.9546e-04
Epoch 44/50
5/5          0s 3ms/step - loss:
6.1022e-04
Epoch 45/50
5/5          0s 3ms/step - loss:
5.4899e-04
Epoch 46/50
5/5          0s 3ms/step - loss:
4.8869e-04
Epoch 47/50
5/5          0s 3ms/step - loss:
4.3266e-04
Epoch 48/50
5/5          0s 3ms/step - loss:
3.8289e-04
Epoch 49/50
5/5          0s 3ms/step - loss:
3.4270e-04
Epoch 50/50
5/5          0s 3ms/step - loss:
3.0429e-04

```

***** Diagnosis of Regularized Linear Model *****

```

Training error: 0.10068223140329917
Validation error: 17.034471617313823

```

***** Diagnosis of Regularized CNN *****

```

Training error: 0.00034364761086180806
Validation error: 7.044750690460205

```

***** Diagnosis of Regularized Decision Tree *****

```

Training error: 0.1841595571737582
Validation error: 12.795367895428631

```

```
[19]: # Sanity checks

# Check if the variables store numeric values
assert isinstance(reg_train_error, float), "custom_reg_train_error must be a numeric value."
assert isinstance(reg_val_error, float), "custom_reg_val_error must be a numeric value."

# Check if the variables store numeric values
assert isinstance(cnn_train_error, float), "cnn_train_error must be a numeric value."
assert isinstance(cnn_val_error, float), "cnn_val_error must be a numeric value."

# Check if the variables store numeric values
assert isinstance(tree_train_error, float), "tree_train_error must be a numeric value."
assert isinstance(tree_val_error, float), "tree_val_error must be a numeric value."

print('Sanity check passed!')
```

Sanity check passed!

1.5 Results

```
[20]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

# Define model names
models = ["Linear Model", "CNN", "Decision Tree"]

# Training errors
train_errors = [reg_train_error, cnn_train_error, tree_train_error]
custom_train_errors = [custom_reg_train_error, custom_cnn_train_error, custom_tree_train_error]

# Validation errors
val_errors = [reg_val_error, cnn_val_error, tree_val_error]
custom_val_errors = [custom_reg_val_error, custom_cnn_val_error, custom_tree_val_error]

# Create a DataFrame for better visualization
error_data = pd.DataFrame({
    "Model": models,
    "Training Error (Original)": train_errors,
```

```

    "Training Error (Regularized)": custom_train_errors,
    "Validation Error (Original)": val_errors,
    "Validation Error (Regularized)": custom_val_errors
})

# Print the table
print(error_data.to_string(index=False))



# Define x-axis positions for grouped bar charts
x = np.arange(len(models)) # Model positions
width = 0.3 # Bar width for better visualization

# Plot Training Errors
plt.figure(figsize=(10, 5))
plt.bar(x - width/2, train_errors, width, label='Original Model', color='blue',
        alpha=0.7)
plt.bar(x + width/2, custom_train_errors, width, label='Regularized Model',
        color='red', alpha=0.7)
plt.xlabel("Models")
plt.ylabel("Training Error")
plt.title("Training Errors (Regularized vs. Non-Regularized)")
plt.xticks(ticks=x, labels=models)
plt.legend()
plt.show()

# Plot Validation Errors
plt.figure(figsize=(10, 5))
plt.bar(x - width/2, val_errors, width, label='Original Model', color='blue',
        alpha=0.7)
plt.bar(x + width/2, custom_val_errors, width, label='Regularized Model',
        color='red', alpha=0.7)
plt.xlabel("Models")
plt.ylabel("Validation Error")
plt.title("Validation Errors (Regularized vs. Non-Regularized)")
plt.xticks(ticks=x, labels=models)
plt.legend()
plt.show()

```

	Model Training Error (Original)	Training Error (Regularized)
Validation Error (Original)		
Linear Model	4.561860e-30	0.100682
22.369070	17.034472	
CNN	4.985445e-03	0.000344
7.242134	7.044751	
Decision Tree	2.381357e-02	0.184160
12.172041	12.795368	



[]: